


DirtyCred: Escalating Privilege in Linux Kernel

Zhenpeng Lin


11/07/2022

How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap



(a) Type confusion between Type A and B




(b) Partial overlap between Type C and A


How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives



(a) Type confusion between Type A and B




(b) Partial overlap between Type C and A

How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap
- Leak kernel pointers
- Tamper kernel pointers

Obtain Primitives



Partial overlap between Type C and A

How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives

- Leak kernel pointers
- Tamper kernel pointers

Bypass Mitigation

How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives

- Leak kernel pointers
- Tamper kernel pointers

Bypass Mitigation

- Execute ROP in different forms^[1]

[1] [Joy of exploiting the kernel](#)

How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives

- Leak kernel pointers
- Tamper kernel pointers

Bypass Mitigation

- Execute ROP in different forms^[1]

Escalate Privilege

[1] [Joy of exploiting the kernel](#)

How Researchers Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives

- Leak kernel pointers
- Tamper kernel pointers

Bypass Mitigation

- Execute ROP in different forms^[1]

Escalate Privilege

Used by 15/17 exploits in [2]

[1] [Joy of exploiting the kernel](#)

[2] [Kernel Exploit Recipes Notebook](#)

How DirtyCred Exploits Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and overlap
- Leak kernel pointers
- Tamper kernel pointers
- Execute ROP

How DirtyCred Exploits Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap
- Leak kernel pointers
- Tamper kernel pointers
- Execute ROP

Obtain Primitives

How DirtyCred Exploits Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives

- ~~Leak kernel pointers~~
- ~~Tamper kernel pointers~~
- ~~Execute ROP~~

How DirtyCred Exploit Kernel Vulns

- Spatial/Temporal memory error
- Type confusion and memory overlap

Obtain Primitives

- Swap kernel credentials

Escalate Privilege

Kernel Credential

- Properties that carry privilege information in kernel
 - Defined in kernel documentation
 - Representation of **privilege** and **capability**
 - Two main types: ***task credentials*** and ***open file credentials***
 - Security checks act on credential objects


Task Credential

- Struct `cred` in Linux kernel's implementation


```
struct cred {  
    atomic_t           usage;  
#ifdef CONFIG_DEBUG_CREDENTIALS  
    atomic_t           subscribers; /* number of processes subscribed */  
    void              *put_addr;  
    unsigned           magic;  
#define CRED_MAGIC      0x43736564  
#define CRED_MAGIC_DEAD 0x44656144  
#endif  
    kuid_t             uid;          /* real UID of the task */  
    kgid_t             gid;          /* real GID of the task */  
    kuid_t             suid;         /* saved UID of the task */  
    kgid_t             sgid;         /* saved GID of the task */  
    kuid_t             euid;         /* effective UID of the task */  
    kgid_t             egid;         /* effective GID of the task */  
    kuid_t             fsuid;        /* UID for VFS ops */  
    kgid_t             fsgid;        /* GID for VFS ops */
```

Task Credential


- **Struct cred** in Linux kernel's implementation
- Represents the *privilege* of kernel tasks



How Linux Kernel Uses Task Credential



How Linux Kernel Uses Task Credential



Open File Credential


- **Struct file** in Linux kernel's implementation

```
struct file {
    union {
        struct llist_node      f_llist;
        struct rcu_head         f_rcuhead;
        unsigned int             f_iocb_flags;
    };
    struct path              f_path;
    struct inode             *f_inode;           /* cache */
    const struct file_operations *f_op;


    /*
     * Protects f_ep, f_flags.
     * Must not be taken from IRQ context.
     */
    spinlock_t               f_lock;
    atomic_long_t             f_count;
    unsigned int               f_flags;
    fmode_t                  f_mode;
    struct mutex              f_pos_lock;
    loff_t                    f_pos;
    struct fown_struct        f_owner;
    const struct cred          *f_cred;
    struct file_ra_state      f_ra;
```

Open File Credential


- Carries the information of opened files (e.g. mode, path, etc.)




How Linux Kernel Uses Open File Credential



How Linux Kernel Uses Open File Credential




Attacking Task Credential




Attacking Task Credential

Step 1. Free the *unprivileged* credential with the vulnerability




Attacking Task Credential

Step 1. Free the *unprivileged* credential with the vulnerability




Attacking Task Credential

Step 2. Allocate a privileged credential in the **freed** memory slot




Attacking Task Credential

Result: Becoming a *privileged* user




Attacking Task Credential

Result: Becoming a *privileged* user




Attacking Open File Credential




Attacking Open File Credential

Step 1. Free a *read-write* file *after* checks, but *before* writing to disk




Attacking Open File Credential

Step 1. Free a *read-write* file *after* checks, but *before* writing to disk




Attacking Open File Credential

Step 2. Allocate a *read-only* file in the **freed** memory slot



Attacking Open File Credential

Result: Writing content to read-only files



Challenges


1. How to **free** credentials.
2. How to **allocate** privileged credentials as unprivileged users.
(attacking *task* credentials)
3. How to finish attack in a **small** time window. (attacking *open file* credentials)

Challenges

1. How to **free** credentials.
2. How to **allocate** *privileged* credentials as *unprivileged* users.
(attacking *task* credentials)
3. How to finish attack in a **small** time window. (attacking *open file* credentials)

Challenge 1: Free Credentials Invalidly

- Both *cred* and *file* object are in **dedicated** caches
- Most vulnerabilities happens in **generic** caches




Challenge 1: Free Credentials Invalidly

- **Solution: Pivoting Vulnerability Capability**
 - Pivoting Invalid-Write (e.g., OOB & UAF write)
 - Pivoting Invalid-Free (e.g., Double-Free)

Pivoting Invalid-Write

Pivoting Invalid-Write

- Leverage victim objects with a reference to credentials




```
struct request_key_auth {
    struct rcu_head
    struct key
    struct key
    const struct cred
    void
    size_t
    pid_t
    char
} __randomize_layout;
```

rcu;
*target_key;
*dest_keyring;
*cred;
*callout_info;
callout_len;
pid;
op[8];


Pivoting Invalid-Write

- Manipulate the memory layout to put the **cred** in the overwrite region




Pivoting Invalid-Write

- *Partially* overwrite the pointer to cause a reference unbalance



Pivoting Invalid-Write


- Free the credential object when freeing the victim object




Pivoting Invalid-Free

Pivoting Invalid-Free

- Two references to free the same object



Pivoting Invalid-Free



**Step 1. Trigger the vuln, free the vuln object
with one reference**


Pivoting Invalid-Free



**Step 1. Trigger the vuln, free the vuln object
with one reference**


**Step 2. Free the object in the memory cache
to free the memory page**

Pivoting Invalid-Free




Step 1. Trigger the vuln, free the vuln object with one reference

Step 2. Free the object in the memory cache to free the memory page




Step 3. Allocate credentials to reclaim the *freed* memory page (*Cross Cache Attack*)


Pivoting Invalid-Free




Step 1. Trigger the vuln, free the vuln object with one reference



Step 2. Free the object in the memory cache to free the memory page



Step 3. Allocate credentials to reclaim the freed memory page (Cross Cache Attack)



Step 4. Free the credentials with the left dangling reference

Challenges

1. How to **free** credentials.
2. How to **allocate privileged credentials as unprivileged users.**
(attacking *task* credentials)
3. How to finish attack in a **small** time window. (attacking *open file* credentials)

Challenge 2: Allocating Privileged Task Credentials

- *Unprivileged* users come with *unprivileged* task credentials
- Waiting privileged users to allocate task credentials influences the success rate

Challenge 2: Allocating Privileged Task Credentials

- **Solution I: Triggering Privileged Userspace Process**
 - Executables with root SUID (e.g. su, mount)
 - Daemons running as root (e.g. sshd)

Challenge 2: Allocating Privileged Task Credentials


- **Solution I: Triggering Privileged Userspace Process**
 - Executables with root SUID (e.g. su, mount)
 - Daemons running as root (e.g. sshd)
- **Solution II: Triggering Privileged Kernel Thread**
 - Kernel Workqueue — spawn new workers
 - Usermode helper — load kernel modules from userspace

Challenges

1. How to **free** credentials.
2. How to **allocate** *privileged* credentials as *unprivileged* users.
(attacking *task* credentials)
3. How to finish attack in a **small** time window. (attacking *open file* credentials)


Challenge 3: Wining the race

- Kernel will examine the access permission before writing to the disk




Challenge 3: Wining the race

- The swap of *file* object happens before permission check




Challenge 3: Wining the race

- The swap of *file* object happens before permission check




Challenge 3: Wining the race

- The swap of *file* object happens before permission check




Challenge 3: Wining the race

- The swap of *file* object happens after *file write*.




Challenge 3: Wining the race

- The swap of *file* object happens after *file write*.




Challenge 3: Wining the race

- The swap of *file* object happens after *file write*.




Challenge 3: Wining the race

- The swap happens in between permission check and file write




Challenge 3: Wining the race

- The swap happens in between permission check and file write




Challenge 3: Wining the race

- The swap happens in between permission check and file write



Challenge 3: Wining the race

- The swap must happen after permission check and before file write



Challenge 3: Wining the race

- **Solution I: Extending with Userfaultfd or FUSE**
 - *Pause kernel execution when accessing userspace memory*

Solution I: Userfaultfd & FUSE

- Pause at *import_iovec* before v4.13
 - *import_iovec* copies userspace memory

```
ssize_t vfs_writev(...)  
{  
    // permission checks  
    if (!(file->f_mode & FMODE_WRITE))  
        return -EBADF;  
    if (!(file->f_mode & FMODE_CAN_WRITE))  
        return -EINVAL;  
  
    ...  
    // import iovec to kernel, where kernel would be paused  
    // using userfaultfd & FUSE  
    res = import_iovec(type, uvector, nr_segs,  
                      ARRAY_SIZE iovstack), &iov, &iter);  
    ...  
    // do file writev  
}
```

Solution I: Userfaultfd & FUSE

- **Pause at *import_iovec* before v4.13**
 - *import_iovec* copies userspace memory
 - Used in Jann Horn's exploitation for [CVE-2016-4557](#)
 - *Dead after v4.13*

Solution I: Userfaultfd & FUSE

- **vfs_writev after v4.13**

```
ssize_t vfs_writev(...)  
{  
    ...  
    // import iovec to kernel, where kernel would be paused  
    // using userfaultfd  
    res = import_iovec(type, uvector, nr_segs,  
                       ARRAY_SIZE(iovstack), &iov, &iter);  
    ...  
    // permission checks  
    if (!(file->f_mode & FMODE_WRITE))  
        return -EBADF;  
    if (!(file->f_mode & FMODE_CAN_WRITE))  
        return -EINVAL;  
    ...  
    // do file writev  
}
```

Solution I: Userfaultfd & FUSE

- Pause at `generic_perform_write`
 - *prefaults* user pages
 - Pauses kernel execution at the page fault

Challenge 3: Wining the race

- **Solution I: Extending with Userfaultfd & FUSE**
 - Pause kernel execution when accessing userspace memory
 - Userfaultfd & FUSE might not be available
- **Solution II: Extending with file lock**
 - Pause kernel execution with lock


Solution II: File Lock

- A lock of the *inode* of the file
- Lock the file when it is being writing to


```
static ssize_t ext4_buffered_write_iter(struct kiocb *iocb,
                                       struct iov_iter *from)
{
    ssize_t ret;
    struct inode *inode = file_inode(iocb->ki_filp);
    inode_lock(inode);
    ...
    ret = generic_perform_write(iocb->ki_filp, from,
                                iocb->ki_pos);
    ...
    inode_unlock(inode);
    return ret;
}
```

Solution II: File Lock

Thread A




Thread B




Solution II: File Lock

Thread A



Thread B




Demo Time!

Northwestern

CVE-2021-4154

Centos 8 and Ubuntu 20




```
[low@centos8 ~]$ █
```

```
low@ubuntu:~$ id █
```

Android Kernel with CFI enabled*

```
oriole:/data/local/tmp $ █
```



Real-World Impact

- **CVE-2021-4154**
 - Received rewards from Google's KCTF
 - The exploit works across kernel v4.18 ~ v5.10
- **CVE-2022-2588**
 - Pwn2own exploitation
 - The exploit works across kernel v3.17 ~ v5.19
- **CVE-2022-20409**
 - Received rewards from Google's KCTF and Android
 - The exploit works on both Android and generic Linux kernel

Advantages of DirtyCred

- **Simple but effective**
 - Shorter exploit chain with fewer steps
- **No effective mitigation**
 - A new exploitation path, can bypass AUTOSLAB
 - No need to deal with KASLR, KCFI, KPTI, SMAP/SMEP
- **Exploitation friendly**
 - Make your exploit **universal!**

Defense Against DirtyCred

- **Fundamental problem**
 - Object isolation is based on *type* not *privilege*
- **Solution**
 - *Isolate privileged* credentials from *unprivileged* ones
- **Where to isolate?**
 - Virtual memory (privileged credentials will be *vmalloc*-ed)

Code is available at <https://github.com/markakd/DirtyCred>

Overhead of The Defense

Benchmark	Vanilla	Hardened	Overhead
Phoronix			
Apache (Reqs/s)	28603.29	29216.48	-2.14%
Sys-RAM (MB/s)	10320.08	10181.91	1.34%
Sys-CPU (Events/s)	4778.41	4776.69	0.04%
FFmpeg(s)	7.456	7.499	0.58%
OpenSSL (Byte/s)	1149941360	1150926390	-0.09%
OpenSSL (Sign/s)	997.2	993.2	0.40%
PHPBench (Score)	571583	571037	0.09%
PyBench (ms)	1303	1311	0.61%
GIMP (s)	12.357	12.347	-0.08%
PostMark (TPS)	5034	5034	0%
LMBench			
Context Switch (ms)	2.60	2.57	-1.15%
UDP (ms)	9.2	9.26	0.65%
TCP (ms)	12.75	12.73	-0.16%
10k File Create (ms)	13.8	14.79	7.17%
10k File Delete (ms)	6.35	6.62	4.25%
Mmap (ms)	80.23	81.91	2.09%
Pipe (MB/s)	4125.3	4028.9	2.34%
AF Unix (MB/s)	8423.5	8396.7	0.32%
TCP (MB/s)	6767.4	6693.3	1.09%
File Reread (MB/s)	8380.43	8380.65	0%
Mmap Reread (MB/s)	15.7K	15.69K	0.06%
Mem Read (MB/s)	10.9K	10.9K	0%
Mem Write (MB/s)	10.76K	10.77K	-0.09%

Takeaways

- A new exploitation concept — DirtyCred
- Principled approaches to different challenges
- A way to produce *Universal* kernel exploits
- Effective defense with negligible overhead

Zhenpeng Lin ([@Markak](https://zplin.me))

<https://zplin.me>

zplin@u.northwestern.edu

